

Assignment 2

May 9, 2021

*You are currently looking at **version 1.0** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.*

1 Assignment 2 - Introduction to NLTK

In part 1 of this assignment you will use nltk to explore the Herman Melville novel Moby Dick. Then in part 2 you will create a spelling recommender function that uses nltk to find words similar to the misspelling.

1.1 Part 1 - Analyzing Moby Dick

In [69]: `import nltk`

```
nltk.download('punkt')
nltk.download('gutenberg')
nltk.download('genesis')
nltk.download('inaugural')
nltk.download('nps_chat')
nltk.download('webtext')
nltk.download('treebank')
nltk.download('averaged_perceptron_tagger')
nltk.download('words')
from nltk.book import *
import pandas as pd
import numpy as np

# If you would like to work with the raw text you can use 'moby_raw'
with open('moby.txt', 'r') as f:
    moby_raw = f.read()

# If you would like to work with the novel in nltk.Text format you can use 'text1'
moby_tokens = nltk.word_tokenize(moby_raw)
text1 = nltk.Text(moby_tokens)
#print(moby_tokens)
```

```

[nltk_data] Downloading package punkt to /home/jovyan/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package gutenber to /home/jovyan/nltk_data...
[nltk_data] Package gutenber is already up-to-date!
[nltk_data] Downloading package genesis to /home/jovyan/nltk_data...
[nltk_data] Package genesis is already up-to-date!
[nltk_data] Downloading package inaugural to /home/jovyan/nltk_data...
[nltk_data] Package inaugural is already up-to-date!
[nltk_data] Downloading package nps_chat to /home/jovyan/nltk_data...
[nltk_data] Package nps_chat is already up-to-date!
[nltk_data] Downloading package webtext to /home/jovyan/nltk_data...
[nltk_data] Package webtext is already up-to-date!
[nltk_data] Downloading package treebank to /home/jovyan/nltk_data...
[nltk_data] Package treebank is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] /home/jovyan/nltk_data...
[nltk_data] Package averaged_perceptron_tagger is already up-to-
[nltk_data] date!
[nltk_data] Downloading package words to /home/jovyan/nltk_data...
[nltk_data] Package words is already up-to-date!

```

1.1.1 Example 1

How many tokens (words and punctuation symbols) are in text1?

This function should return an integer.

```

In [70]: def example_one():

           return len(nltk.word_tokenize(moby_raw)) # or alternatively len(text1)

example_one()

```

Out[70]: 254989

1.1.2 Example 2

How many unique tokens (unique words and punctuation) does text1 have?

This function should return an integer.

```

In [71]: def example_two():

           return len(set(nltk.word_tokenize(moby_raw))) # or alternatively len(set(text1))

example_two()

```

Out[71]: 20755

1.1.3 Example 3

After lemmatizing the verbs, how many unique tokens does text1 have?

This function should return an integer.

```
In [72]: from nltk.stem import WordNetLemmatizer
```

```
def example_three():  
  
    lemmatizer = WordNetLemmatizer()  
    lemmatized = [lemmatizer.lemmatize(w, 'v') for w in text1]  
  
    return len(set(lemmatized))  
  
example_three()
```

```
-----  
LookupError                                Traceback (most recent call last)
```

```
/opt/conda/lib/python3.6/site-packages/nltk/corpus/util.py in __load(self)  
79         except LookupError as e:  
--> 80             try: root = nltk.data.find('{}/{}'.format(self.subdir, zip_name))  
81                 except LookupError: raise e
```

```
/opt/conda/lib/python3.6/site-packages/nltk/data.py in find(resource_name, paths)  
674     resource_not_found = '\n%s\n%s\n%s\n' % (sep, msg, sep)  
--> 675     raise LookupError(resource_not_found)  
676
```

```
LookupError:
```

```
*****
```

```
Resource wordnet not found.
```

```
Please use the NLTK Downloader to obtain the resource:
```

```
>>> import nltk  
>>> nltk.download('wordnet')
```

```
Searched in:
```

- '/home/jovyan/nltk_data'
- '/usr/share/nltk_data'
- '/usr/local/share/nltk_data'
- '/usr/lib/nltk_data'
- '/usr/local/lib/nltk_data'
- '/opt/conda/nltk_data'
- '/opt/conda/share/nltk_data'

```
- '/opt/conda/lib/nltk_data'
*****
```

During handling of the above exception, another exception occurred:

```
LookupError                                Traceback (most recent call last)

<ipython-input-72-6097938ac8b3> in <module>()
      8     return len(set(lemmatized))
      9
----> 10 example_three()

<ipython-input-72-6097938ac8b3> in example_three()
      4
      5     lemmatizer = WordNetLemmatizer()
----> 6     lemmatized = [lemmatizer.lemmatize(w,'v') for w in text1]
      7
      8     return len(set(lemmatized))

<ipython-input-72-6097938ac8b3> in <listcomp>(.0)
      4
      5     lemmatizer = WordNetLemmatizer()
----> 6     lemmatized = [lemmatizer.lemmatize(w,'v') for w in text1]
      7
      8     return len(set(lemmatized))

/opt/conda/lib/python3.6/site-packages/nltk/stem/wordnet.py in lemmatize(self, word, pos)
    38
    39     def lemmatize(self, word, pos=NOUN):
----> 40         lemmas = wordnet._morphology(word, pos)
    41         return min(lemmas, key=len) if lemmas else word
    42

/opt/conda/lib/python3.6/site-packages/nltk/corpus/util.py in __getattr__(self, attr)
    114         raise AttributeError("LazyCorpusLoader object has no attribute '__bases__'")
    115
--> 116         self.__load()
    117         # This looks circular, but its not, since __load() changes our
    118         # __class__ to something new:
```

```

/opt/conda/lib/python3.6/site-packages/nltk/corpus/util.py in __load(self)
    79         except LookupError as e:
    80             try: root = nltk.data.find('{}/{}'.format(self.subdir, zip_name))
--> 81             except LookupError: raise e
    82
    83         # Load the corpus.

```

```

/opt/conda/lib/python3.6/site-packages/nltk/corpus/util.py in __load(self)
    76         else:
    77             try:
--> 78                 root = nltk.data.find('{}/{}'.format(self.subdir, self.__name))
    79             except LookupError as e:
    80                 try: root = nltk.data.find('{}/{}'.format(self.subdir, zip_name))

```

```

/opt/conda/lib/python3.6/site-packages/nltk/data.py in find(resource_name, paths)
    673     sep = '*' * 70
    674     resource_not_found = '\n%s\n%s\n%s\n' % (sep, msg, sep)
--> 675     raise LookupError(resource_not_found)
    676
    677

```

```

LookupError:
*****
Resource wordnet not found.
Please use the NLTK Downloader to obtain the resource:

```

```

>>> import nltk
>>> nltk.download('wordnet')

```

```

Searched in:
- '/home/jovyan/nltk_data'
- '/usr/share/nltk_data'
- '/usr/local/share/nltk_data'
- '/usr/lib/nltk_data'
- '/usr/local/lib/nltk_data'
- '/opt/conda/nltk_data'
- '/opt/conda/share/nltk_data'
- '/opt/conda/lib/nltk_data'

```

```

*****

```

1.1.4 Question 1

What is the lexical diversity of the given text input? (i.e. ratio of unique tokens to the total number of tokens)

This function should return a float.

```
In [ ]: def answer_one():

    return float(example_two()/example_one())# Your answer here

answer_one()
```

1.1.5 Question 2

What percentage of tokens is 'whale' or 'Whale'?

This function should return a float.

```
In [ ]: def answer_two():

    text = nltk.word_tokenize(moby_raw)
    count = 0
    for t in text:
        if t == 'whale' or t == 'Whale':
            count+=1
    return 100 * float(count/example_one())# Your answer here

answer_two()
```

1.1.6 Question 3

What are the 20 most frequently occurring (unique) tokens in the text? What is their frequency?

This function should return a list of 20 tuples where each tuple is of the form (token, frequency). The list should be sorted in descending order of frequency.

```
In [ ]: def answer_three():

    dist = FreqDist(moby_tokens)
    import operator
    sorted_d = dict(sorted(dist.items(), key=operator.itemgetter(1), reverse=True))
    list_1 = []
    count = 0
    for k, v in sorted_d.items():
        if count < 20:
            list_1.append((k,v))
            count+=1
    return list_1 # Your answer here

answer_three()
```

1.1.7 Question 4

What tokens have a length of greater than 5 and frequency of more than 150?

This function should return an alphabetically sorted list of the tokens that match the above constraints. To sort your list, use `sorted()`

```
In [ ]: def answer_four():

    dist = FreqDist(moby_tokens)
    words = dist.keys()
    my_words = [w for w in words if len(w)>5 and dist[w]>150]
    return sorted(my_words) # Your answer here

answer_four()
```

1.1.8 Question 5

Find the longest word in text1 and that word's length.

This function should return a tuple (longest_word, length).

```
In [ ]: def answer_five():
    length1 = 0
    longest = ''
    for w in text1:
        if len(w) > length1:
            length1 = len(w)
            longest = w

    return longest, length1 # Your answer here

answer_five()
```

1.1.9 Question 6

What unique words have a frequency of more than 2000? What is their frequency?

"Hint: you may want to use `isalpha()` to check if the token is a word and not punctuation."

This function should return a list of tuples of the form (frequency, word) sorted in descending order of frequency.

```
In [ ]: def answer_six():

    dist = FreqDist(moby_tokens)
    words = dist.keys()
    words_1 = [w for w in words if w.isalpha()]
    my_words = [w for w in words_1 if dist[w]>2000]
    my_list = []
    for w in my_words:
        my_list.append((dist[w],w))
```

```

final_list = sorted(my_list, reverse=True)

return final_list# Your answer here

answer_six()

```

1.1.10 Question 7

What is the average number of tokens per sentence?

This function should return a float.

```

In [ ]: def answer_seven():

    sentences = nltk.sent_tokenize(moby_raw)
    return float(len(text1)/len(sentences)) #Your answer here

answer_seven()

```

1.1.11 Question 8

What are the 5 most frequent parts of speech in this text? What is their frequency?

This function should return a list of tuples of the form (part_of_speech, frequency) sorted in descending order of frequency.

```

In [98]: def answer_eight():

    tags = nltk.pos_tag(moby_tokens)
    dist = FreqDist([tag for (word, tag) in tags])
    #final_list = sorted(dist, reverse=True)

    return dist.most_common(5)# Your answer here

# Your answer here

answer_eight()

Out[98]: [('NN', 32730), ('IN', 28657), ('DT', 25867), ('.', 19204), ('JJ', 17620)]

```

1.2 Part 2 - Spelling Recommender

For this part of the assignment you will create three different spelling recommenders, that each take a list of misspelled words and recommends a correctly spelled word for every word in the list.

For every misspelled word, the recommender should find the word in `correct_spellings` that has the shortest distance*, and starts with the same letter as the misspelled word, and return that word as a recommendation.

*Each of the three different recommenders will use a different distance measure (outlined below).

Each of the recommenders should provide recommendations for the three default words provided: ['cormulent', 'incendenece', 'validate'].


```
In [ ]: from nltk.corpus import words
        from nltk.metrics.distance import (
            edit_distance,
            jaccard_distance,
        )
        from nltk.util import ngrams

        correct_spellings = words.words()
```

1.2.1 Question 9

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Jaccard distance on the trigrams of the two words.

This function should return a list of length three: ['cormulent_reccomendation', 'incendenece_reccomendation', 'validate_reccomendation'].

```
In [75]: spellings_series = pd.Series(correct_spellings)
        def jaccard(entries, gram_number):
            outcomes = []
            for entry in entries:
                spellings = spellings_series[spellings_series.str.startswith(entry[0])]
                distances = ((jaccard_distance(set(ngrams(entry, gram_number)),
                                                    set(ngrams(word, gram_number))), word)
                             for word in spellings)
                closest = min(distances)
                outcomes.append(closest[1])
            return outcomes
        def answer_nine(entries=['cormulent', 'incendenece', 'validate']):
            #finds the closest word based on jaccard distance
            return jaccard(entries, 3)

        #print(answer_nine())

        answer_nine()
```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:17: DeprecationWarning: generator 'generator' is deprecated

```
Out[75]: ['corpulent', 'indecence', 'validate']
```

1.2.2 Question 10

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Jaccard distance on the 4-grams of the two words.

This function should return a list of length three: ['cormulent_reccomendation', 'incendenece_reccomendation', 'validate_reccomendation'].

```
In [76]: def answer_ten(entries=['cormulent', 'incendenece', 'validate']):
```

```
    return jaccard(entries, 4)
```

```
    answer_ten()
```

```
/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:17: DeprecationWarning: generator ' 
```

```
Out[76]: ['cormus', 'incendiary', 'valid']
```

1.2.3 Question 11

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Edit distance on the two words with transpositions.

This function should return a list of length three: ['cormulent_reccomendation', 'incendenece_reccomendation', 'validate_reccomendation'].

```
In [77]: def answer_eleven(entries=['cormulent', 'incendenece', 'validate']):
```

```
    outcomes = []
```

```
    for entry in entries:
```

```
        distances = ((edit_distance(entry,
                                   word), word)
```

```
                     for word in correct_spellings)
```

```
        closest = min(distances)
```

```
        outcomes.append(closest[1])
```

```
    return outcomes
```

```
    answer_eleven()
```

```
Out[77]: ['corpulent', 'intendence', 'validate']
```

```
In [ ]:
```