

PandasIdioms_ed

March 1, 2021

Python programmers will often suggest that there many ways the language can be used to solve a particular problem. But that some are more appropriate than others. The best solutions are celebrated as Idiomatic Python and there are lots of great examples of this on StackOverflow and other websites.

A sort of sub-language within Python, Pandas has its own set of idioms. We've alluded to some of these already, such as using vectorization whenever possible, and not using iterative loops if you don't need to. Several developers and users within the Panda's community have used the term **pandorable** for these idioms. I think it's a great term. So, I wanted to share with you a couple of key features of how you can make your code pandorable.

```
[1]: # Let's start by bringing in our data processing libraries
import pandas as pd
import numpy as np
# And we'll bring in some timing functionality too, from the timeit module
import timeit

# And lets look at some census data from the US
df = pd.read_csv('datasets/census.csv')
df.head()
```

```
[1]:
```

	SUMLEV	REGION	DIVISION	STATE	COUNTY	STNAME	CTYNAME	\
0	40	3	6	1	0	Alabama	Alabama	
1	50	3	6	1	1	Alabama	Autauga County	
2	50	3	6	1	3	Alabama	Baldwin County	
3	50	3	6	1	5	Alabama	Barbour County	
4	50	3	6	1	7	Alabama	Bibb County	

	CENSUS2010POP	ESTIMATESBASE2010	POPESTIMATE2010	...	RDOMESTICMIG2011	\
0	4779736	4780127	4785161	...	0.002295	
1	54571	54571	54660	...	7.242091	
2	182265	182265	183193	...	14.832960	
3	27457	27457	27341	...	-4.728132	
4	22915	22919	22861	...	-5.527043	

	RDOMESTICMIG2012	RDOMESTICMIG2013	RDOMESTICMIG2014	RDOMESTICMIG2015	\
0	-0.193196	0.381066	0.582002	-0.467369	
1	-2.915927	-3.012349	2.265971	-2.530799	
2	17.647293	21.845705	19.243287	17.197872	

3	-2.500690	-7.056824	-3.904217	-10.543299
4	-5.068871	-6.201001	-0.177537	0.177258

	RNETMIG2011	RNETMIG2012	RNETMIG2013	RNETMIG2014	RNETMIG2015
0	1.030015	0.826644	1.383282	1.724718	0.712594
1	7.606016	-2.626146	-2.722002	2.592270	-2.187333
2	15.844176	18.559627	22.727626	20.317142	18.293499
3	-4.874741	-2.758113	-7.167664	-3.978583	-10.543299
4	-5.088389	-4.363636	-5.403729	0.754533	1.107861

[5 rows x 100 columns]

```
[2]: # The first of the pandas idioms I would like to talk about is called method chaining. The general idea behind
# method chaining is that every method on an object returns a reference to that object. The beauty of this is
# that you can condense many different operations on a DataFrame, for instance, into one line or at least one
# statement of code.

# Here's the pandorable way to write code with method chaining. In this code I'm going to pull out the state
# and city names as a multiple index, and I'm going to do so only for data which has a summary level of 50,
# which in this dataset is county-level data. I'll rename a column too, just to make it a bit more readable.
(df.where(df['SUMLEV']==50)
 .dropna()
 .set_index(['STNAME','CTYNAME'])
 .rename(columns={'ESTIMATESBASE2010': 'Estimates Base 2010'}))
```

```
[2]:
```

		SUMLEV	REGION	DIVISION	STATE	COUNTY	\
	STNAME CTYNAME						
	Alabama Autauga County	50.0	3.0	6.0	1.0	1.0	
	Baldwin County	50.0	3.0	6.0	1.0	3.0	
	Barbour County	50.0	3.0	6.0	1.0	5.0	
	Bibb County	50.0	3.0	6.0	1.0	7.0	
	Blount County	50.0	3.0	6.0	1.0	9.0	
	
	Wyoming Sweetwater County	50.0	4.0	8.0	56.0	37.0	
	Teton County	50.0	4.0	8.0	56.0	39.0	
	Uinta County	50.0	4.0	8.0	56.0	41.0	
	Washakie County	50.0	4.0	8.0	56.0	43.0	
	Weston County	50.0	4.0	8.0	56.0	45.0	

		CENSUS2010POP	Estimates Base 2010	\
	STNAME CTYNAME			

Alabama	Autauga County	54571.0	54571.0
	Baldwin County	182265.0	182265.0
	Barbour County	27457.0	27457.0
	Bibb County	22915.0	22919.0
	Blount County	57322.0	57322.0
...
Wyoming	Sweetwater County	43806.0	43806.0
	Teton County	21294.0	21294.0
	Uinta County	21118.0	21118.0
	Washakie County	8533.0	8533.0
	Weston County	7208.0	7208.0

STNAME	CTYNAME	POPESTIMATE2010	POPESTIMATE2011	POPESTIMATE2012	\
Alabama	Autauga County	54660.0	55253.0	55175.0	
	Baldwin County	183193.0	186659.0	190396.0	
	Barbour County	27341.0	27226.0	27159.0	
	Bibb County	22861.0	22733.0	22642.0	
	Blount County	57373.0	57711.0	57776.0	
...	
Wyoming	Sweetwater County	43593.0	44041.0	45104.0	
	Teton County	21297.0	21482.0	21697.0	
	Uinta County	21102.0	20912.0	20989.0	
	Washakie County	8545.0	8469.0	8443.0	
	Weston County	7181.0	7114.0	7065.0	

STNAME	CTYNAME	...	RDOMESTICMIG2011	RDOMESTICMIG2012	\
Alabama	Autauga County	...	7.242091	-2.915927	
	Baldwin County	...	14.832960	17.647293	
	Barbour County	...	-4.728132	-2.500690	
	Bibb County	...	-5.527043	-5.068871	
	Blount County	...	1.807375	-1.177622	
...	
Wyoming	Sweetwater County	...	1.072643	16.243199	
	Teton County	...	-1.589565	0.972695	
	Uinta County	...	-17.755986	-4.916350	
	Washakie County	...	-11.637475	-0.827815	
	Weston County	...	-11.752361	-8.040059	

STNAME	CTYNAME	RDOMESTICMIG2013	RDOMESTICMIG2014	\
Alabama	Autauga County	-3.012349	2.265971	
	Baldwin County	21.845705	19.243287	
	Barbour County	-7.056824	-3.904217	
	Bibb County	-6.201001	-0.177537	
	Blount County	-1.748766	-2.062535	

```

...
Wyoming Sweetwater County      -5.339774      -14.252889
        Teton County           19.525929       14.143021
        Uinta County           -6.902954      -14.215862
        Washakie County        -2.013502      -17.781491
        Weston County          12.372583       1.533635

                                RDOMESTICMIG2015  RNETMIG2011  RNETMIG2012  \
STNAME CTYNAME
Alabama Autauga County          -2.530799       7.606016     -2.626146
        Baldwin County         17.197872      15.844176     18.559627
        Barbour County        -10.543299      -4.874741     -2.758113
        Bibb County           0.177258      -5.088389     -4.363636
        Blount County         -1.369970       1.859511     -0.848580
...
Wyoming Sweetwater County      -14.248864       1.255221     16.243199
        Teton County          -0.564849       0.654527       2.408578
        Uinta County          -12.127022      -18.136812     -5.536861
        Washakie County        1.682288      -11.990126     -1.182592
        Weston County          6.935294      -12.032179     -8.040059

                                RNETMIG2013  RNETMIG2014  RNETMIG2015
STNAME CTYNAME
Alabama Autauga County          -2.722002       2.592270     -2.187333
        Baldwin County         22.727626      20.317142     18.293499
        Barbour County        -7.167664      -3.978583    -10.543299
        Bibb County           -5.403729       0.754533       1.107861
        Blount County         -1.402476      -1.577232     -0.884411
...
Wyoming Sweetwater County      -5.295460     -14.075283    -14.070195
        Teton County          21.160658      16.308671       1.520747
        Uinta County          -7.521840     -14.740608    -12.606351
        Washakie County       -2.250385     -18.020168       1.441961
        Weston County          12.372583       1.533635       6.935294

```

[3142 rows x 98 columns]

```

[3]: # Lets walk through this. First, we use the where() function on the dataframe
      ↳and pass in a boolean mask which
      # is only true for those rows where the SUMLEV is equal to 50. This indicates
      ↳in our source data that the data
      # is summarized at the county level. With the result of the where() function
      ↳evaluated, we drop missing
      # values. Remember that .where() doesn't drop missing values by default. Then
      ↳we set an index on the result of
      # that. In this case I've set it to the state name followed by the county name.
      ↳Finally. I rename a column to

```

```
# make it more readable. Note that instead of writing this all on one line, as
→I could have done, I began the
# statement with a parenthesis, which tells python I'm going to span the
→statement over multiple lines for
# readability.
```

```
[4]: # Here's a more traditional, non-pandorable way, of writing this. There's
→nothing wrong with this code in the
# functional sense, you might even be able to understand it better as a new
→person to the language. It's just
# not as pandorable as the first example.

# First create a new dataframe from the original
df = df[df['SUMLEV']==50] # I'll use the overloaded indexing operator [] which
→drops nans
# Update the dataframe to have a new index, we use inplace=True to do this in
→place
df.set_index(['STNAME', 'CTYNAME'], inplace=True)
# Set the column names
df.rename(columns={'ESTIMATESBASE2010': 'Estimates Base 2010'})
```

```
[4]:
```

		SUMLEV	REGION	DIVISION	STATE	COUNTY	\
STNAME	CTYNAME						
Alabama	Autauga County	50	3	6	1	1	
	Baldwin County	50	3	6	1	3	
	Barbour County	50	3	6	1	5	
	Bibb County	50	3	6	1	7	
	Blount County	50	3	6	1	9	
...	
Wyoming	Sweetwater County	50	4	8	56	37	
	Teton County	50	4	8	56	39	
	Uinta County	50	4	8	56	41	
	Washakie County	50	4	8	56	43	
	Weston County	50	4	8	56	45	

		CENSUS2010POP	Estimates Base 2010	\
STNAME	CTYNAME			
Alabama	Autauga County	54571	54571	
	Baldwin County	182265	182265	
	Barbour County	27457	27457	
	Bibb County	22915	22919	
	Blount County	57322	57322	
...	
Wyoming	Sweetwater County	43806	43806	
	Teton County	21294	21294	
	Uinta County	21118	21118	
	Washakie County	8533	8533	

	Weston County	7208	7208	
		POPESTIMATE2010	POPESTIMATE2011	POPESTIMATE2012 \
STNAME	CTYNAME			
Alabama	Autauga County	54660	55253	55175
	Baldwin County	183193	186659	190396
	Barbour County	27341	27226	27159
	Bibb County	22861	22733	22642
	Blount County	57373	57711	57776
...
Wyoming	Sweetwater County	43593	44041	45104
	Teton County	21297	21482	21697
	Uinta County	21102	20912	20989
	Washakie County	8545	8469	8443
	Weston County	7181	7114	7065
		...	RDOMESTICMIG2011	RDOMESTICMIG2012 \
STNAME	CTYNAME	...		
Alabama	Autauga County	...	7.242091	-2.915927
	Baldwin County	...	14.832960	17.647293
	Barbour County	...	-4.728132	-2.500690
	Bibb County	...	-5.527043	-5.068871
	Blount County	...	1.807375	-1.177622
...
Wyoming	Sweetwater County	...	1.072643	16.243199
	Teton County	...	-1.589565	0.972695
	Uinta County	...	-17.755986	-4.916350
	Washakie County	...	-11.637475	-0.827815
	Weston County	...	-11.752361	-8.040059
		RDOMESTICMIG2013	RDOMESTICMIG2014 \	
STNAME	CTYNAME			
Alabama	Autauga County	-3.012349	2.265971	
	Baldwin County	21.845705	19.243287	
	Barbour County	-7.056824	-3.904217	
	Bibb County	-6.201001	-0.177537	
	Blount County	-1.748766	-2.062535	
...	
Wyoming	Sweetwater County	-5.339774	-14.252889	
	Teton County	19.525929	14.143021	
	Uinta County	-6.902954	-14.215862	
	Washakie County	-2.013502	-17.781491	
	Weston County	12.372583	1.533635	
		RDOMESTICMIG2015	RNETMIG2011	RNETMIG2012 \
STNAME	CTYNAME			
Alabama	Autauga County	-2.530799	7.606016	-2.626146

	Baldwin County	17.197872	15.844176	18.559627
	Barbour County	-10.543299	-4.874741	-2.758113
	Bibb County	0.177258	-5.088389	-4.363636
	Blount County	-1.369970	1.859511	-0.848580
...	
Wyoming	Sweetwater County	-14.248864	1.255221	16.243199
	Teton County	-0.564849	0.654527	2.408578
	Uinta County	-12.127022	-18.136812	-5.536861
	Washakie County	1.682288	-11.990126	-1.182592
	Weston County	6.935294	-12.032179	-8.040059

		RNETMIG2013	RNETMIG2014	RNETMIG2015
STNAME	CTYNAME			
Alabama	Autauga County	-2.722002	2.592270	-2.187333
	Baldwin County	22.727626	20.317142	18.293499
	Barbour County	-7.167664	-3.978583	-10.543299
	Bibb County	-5.403729	0.754533	1.107861
	Blount County	-1.402476	-1.577232	-0.884411
...	
Wyoming	Sweetwater County	-5.295460	-14.075283	-14.070195
	Teton County	21.160658	16.308671	1.520747
	Uinta County	-7.521840	-14.740608	-12.606351
	Washakie County	-2.250385	-18.020168	1.441961
	Weston County	12.372583	1.533635	6.935294

[3142 rows x 98 columns]

```
[5]: # Now, the key with any good idiom is to understand when it isn't helping you.
      # In this case, you can actually
      # time both methods and see which one runs faster

      # We can put the approach into a function and pass the function into the timeit
      # function to count the time the
      # parameter number allows us to choose how many times we want to run the
      # function. Here we will just set it to
      # 10

      # Lets write a wrapper for our first function
      def first_approach():
          global df
          # And we'll just paste our code right here
          return (df.where(df['SUMLEV']==50)
                  .dropna()
                  .set_index(['STNAME', 'CTYNAME'])
                  .rename(columns={'ESTIMATESBASE2010': 'Estimates Base 2010'}))

      # Read in our dataset anew
```

```
df = pd.read_csv('datasets/census.csv')

# And now lets run it
timeit.timeit(first_approach, number=10)
```

[5]: 1.1084833510685712

```
[6]: # Now let's test the second approach. As you may notice, we use our global
      ↪variable df in the function.
      # However, changing a global variable inside a function will modify the
      ↪variable even in a global scope and we
      # do not want that to happen in this case. Therefore, for selecting summary
      ↪levels of 50 only, I create a new
      # dataframe for those records

      # Let's run this for once and see how fast it is
      def second_approach():
          global df
          new_df = df[df['SUMLEV']==50]
          new_df.set_index(['STNAME', 'CTYNAME'], inplace=True)
          return new_df.rename(columns={'ESTIMATESBASE2010': 'Estimates Base 2010'})

      # Read in our dataset anew
      df = pd.read_csv('datasets/census.csv')

      # And now lets run it
      timeit.timeit(second_approach, number=10)
```

[6]: 0.10386669298168272

```
[7]: # As you can see, the second approach is much faster! So, this is a particular
      ↪example of a classic time
      # readability trade off.

      # You'll see lots of examples on stack overflow and in documentation of people
      ↪using method chaining in their
      # pandas. And so, I think being able to read and understand the syntax is
      ↪really worth your time. But keep in
      # mind that following what appears to be stylistic idioms might have
      ↪performance issues that you need to
      # consider as well.
```

```
[8]: # Here's another pandas idiom. Python has a wonderful function called map,
      ↪which is sort of a basis for
      # functional programming in the language. When you want to use map in Python,
      ↪you pass it some function you
      # want called, and some iterable, like a list, that you want the function to be
      ↪applied to. The results are
```



```
# that the function is called against each item in the list, and there's a
→resulting list of all of the
# evaluations of that function.

# Pandas has a similar function called applymap. In applymap, you provide some
→function which should operate
# on each cell of a DataFrame, and the return set is itself a DataFrame. Now I
→think applymap is fine, but I
# actually rarely use it. Instead, I find myself often wanting to map across
→all of the rows in a DataFrame.
# And pandas has a function that I use heavily there, called apply. Let's look
→at an example.
```

[9]: # Let's take a look at our census DataFrame. In this DataFrame, we have five
→columns for population estimates,
with each column corresponding with one year of estimates. It's quite
→reasonable to want to create some new
columns for minimum or maximum values, and the apply function is an easy way
→to do this.

```
# First, we need to write a function which takes in a particular row of data,
→finds a minimum and maximum
# values, and returns a new row of data and returns a new row of data. We'll
→call this function min_max, this
# is pretty straight forward. We can create some small slice of a row by
→projecting the population columns.
# Then use the NumPy min and max functions, and create a new series with a
→label values represent the new
# values we want to apply.
```

```
def min_max(row):
    data = row[['POPESTIMATE2010',
                'POPESTIMATE2011',
                'POPESTIMATE2012',
                'POPESTIMATE2013',
                'POPESTIMATE2014',
                'POPESTIMATE2015']]
    return pd.Series({'min': np.min(data), 'max': np.max(data)})
```

[10]: # Then we just need to call apply on the DataFrame.

```
# Apply takes the function and the axis on which to operate as parameters. Now,
→we have to be a bit careful,
# we've talked about axis zero being the rows of the DataFrame in the past. But
→this parameter is really the
# parameter of the index to use. So, to apply across all rows, which is
→applying on all columns, you pass axis
```

```
# equal to 'columns'.
df.apply(min_max, axis='columns').head()
```

```
[10]:      min      max
0  4785161  4858979
1    54660   55347
2   183193  203709
3    26489   27341
4    22512   22861
```

```
[11]: # Of course there's no need to limit yourself to returning a new series object.
      → If you're doing this as part
      # of data cleaning your likely to find yourself wanting to add new data to the
      → existing DataFrame. In that
      # case you just take the row values and add in new columns indicating the max
      → and minimum scores. This is a
      # regular part of my workflow when bringing in data and building summary or
      → descriptive statistics, and is
      # often used heavily with the merging of DataFrames.
```

```
[12]: # Here's an example where we have a revised version of the function min_max
      → Instead of returning a separate
      # series to display the min and max we add two new columns in the original
      → dataframe to store min and max
```

```
def min_max(row):
    data = row[['POPESTIMATE2010',
                 'POPESTIMATE2011',
                 'POPESTIMATE2012',
                 'POPESTIMATE2013',
                 'POPESTIMATE2014',
                 'POPESTIMATE2015']]
    # Create a new entry for max
    row['max'] = np.max(data)
    # Create a new entry for min
    row['min'] = np.min(data)
    return row
# Now just apply the function across the dataframe
df.apply(min_max, axis='columns')
```

```
[12]:      SUMLEV  REGION  DIVISION  STATE  COUNTY  STNAME  CTYNAME \
0         40      3         6      1      0  Alabama  Alabama
1         50      3         6      1      1  Alabama  Autauga County
2         50      3         6      1      3  Alabama  Baldwin County
3         50      3         6      1      5  Alabama  Barbour County
4         50      3         6      1      7  Alabama  Bibb County
...      ...      ...      ...      ...      ...      ...      ...
3188      50      4         8     56     37  Wyoming  Sweetwater County
```

3189	50	4	8	56	39	Wyoming	Teton County
3190	50	4	8	56	41	Wyoming	Uinta County
3191	50	4	8	56	43	Wyoming	Washakie County
3192	50	4	8	56	45	Wyoming	Weston County

	CENSUS2010POP	ESTIMATESBASE2010	POPESTIMATE2010	...	\
0	4779736	4780127	4785161	...	
1	54571	54571	54660	...	
2	182265	182265	183193	...	
3	27457	27457	27341	...	
4	22915	22919	22861	...	
...	
3188	43806	43806	43593	...	
3189	21294	21294	21297	...	
3190	21118	21118	21102	...	
3191	8533	8533	8545	...	
3192	7208	7208	7181	...	

	RDOMESTICMIG2013	RDOMESTICMIG2014	RDOMESTICMIG2015	RNETMIG2011	\
0	0.381066	0.582002	-0.467369	1.030015	
1	-3.012349	2.265971	-2.530799	7.606016	
2	21.845705	19.243287	17.197872	15.844176	
3	-7.056824	-3.904217	-10.543299	-4.874741	
4	-6.201001	-0.177537	0.177258	-5.088389	
...	
3188	-5.339774	-14.252889	-14.248864	1.255221	
3189	19.525929	14.143021	-0.564849	0.654527	
3190	-6.902954	-14.215862	-12.127022	-18.136812	
3191	-2.013502	-17.781491	1.682288	-11.990126	
3192	12.372583	1.533635	6.935294	-12.032179	

	RNETMIG2012	RNETMIG2013	RNETMIG2014	RNETMIG2015	max	min
0	0.826644	1.383282	1.724718	0.712594	4858979	4785161
1	-2.626146	-2.722002	2.592270	-2.187333	55347	54660
2	18.559627	22.727626	20.317142	18.293499	203709	183193
3	-2.758113	-7.167664	-3.978583	-10.543299	27341	26489
4	-4.363636	-5.403729	0.754533	1.107861	22861	22512
...
3188	16.243199	-5.295460	-14.075283	-14.070195	45162	43593
3189	2.408578	21.160658	16.308671	1.520747	23125	21297
3190	-5.536861	-7.521840	-14.740608	-12.606351	21102	20822
3191	-1.182592	-2.250385	-18.020168	1.441961	8545	8316
3192	-8.040059	12.372583	1.533635	6.935294	7234	7065

[3193 rows x 102 columns]

```
[13]: # Apply is an extremely important tool in your toolkit. The reason I introduced
      → apply here is because you
      # rarely see it used with large function definitions, like we did. Instead, you
      → typically see it used with
      # lambdas. To get the most of the discussions you'll see online, you're going
      → to need to know how to at least
      # read lambdas.

      # Here's You can imagine how you might chain several apply calls with lambdas
      → together to create a readable
      # yet succinct data manipulation script. One line example of how you might
      → calculate the max of the columns
      # using the apply function.
      rows = ['POPESTIMATE2010', 'POPESTIMATE2011', 'POPESTIMATE2012',
      → 'POPESTIMATE2013', 'POPESTIMATE2014',
      → 'POPESTIMATE2015']
      # Now we'll just apply this across the dataframe with a lambda
      df.apply(lambda x: np.max(x[rows]), axis=1).head()
```

```
[13]: 0    4858979
      1     55347
      2    203709
      3     27341
      4     22861
      dtype: int64
```

```
[14]: # If you don't remember lambdas just pause the video for a moment and look up
      → the syntax. A lambda is just an
      # unnamed function in python, in this case it takes a single parameter, x, and
      → returns a single value, in this
      # case the maximum over all columns associated with row x.
```

```
[15]: # The beauty of the apply function is that it allows flexibility in doing
      → whatever manipulation that you
      # desire, as the function you pass into apply can be any customized however you
      → want. Let's say we want to
      # divide the states into four categories: Northeast, Midwest, South, and West
      → We can write a customized
      # function that returns the region based on the state the state regions
      → information is obtained from Wikipedia

      def get_state_region(x):
          northeast = ['Connecticut', 'Maine', 'Massachusetts', 'New Hampshire',
          → 'Rhode Island', 'Vermont', 'New York', 'New
          → Jersey', 'Pennsylvania']
          midwest = ['Illinois', 'Indiana', 'Michigan', 'Ohio', 'Wisconsin', 'Iowa',
          → 'Kansas', 'Minnesota', 'Missouri', 'Nebraska', 'North Dakota',
```

```

        'South Dakota']
south = ['Delaware', 'Florida', 'Georgia', 'Maryland', 'North Carolina',
        'South Carolina', 'Virginia', 'District of Columbia', 'West Virginia',
        'Alabama', 'Kentucky', 'Mississippi', 'Tennessee', 'Arkansas',
        'Louisiana', 'Oklahoma', 'Texas']
west = ['Arizona', 'Colorado', 'Idaho', 'Montana', 'Nevada', 'New Mexico', 'Utah',
        'Wyoming', 'Alaska', 'California', 'Hawaii', 'Oregon', 'Washington']

if x in northeast:
    return "Northeast"
elif x in midwest:
    return "Midwest"
elif x in south:
    return "South"
else:
    return "West"

```

```

[16]: # Now we have the customized function, let's say we want to create a new column
      ↪ called Region, which shows the
      # state's region, we can use the customized function and the apply function to
      ↪ do so. The customized function
      # is supposed to work on the state name column STNAME. So we will set the apply
      ↪ function on the state name
      # column and pass the customized function into the apply function
df['state_region'] = df['STNAME'].apply(lambda x: get_state_region(x))

```

```

[17]: # Now let's see the results
df[['STNAME', 'state_region']].head()

```

```

[17]:  STNAME state_region
0  Alabama          South
1  Alabama          South
2  Alabama          South
3  Alabama          South
4  Alabama          South

```

So there are a couple of Pandas idioms. But I think there's many more, and I haven't talked about them here. So here's an unofficial assignment for you. Go look at some of the top ranked questions on pandas on Stack Overflow, and look at how some of the more experienced authors, answer those questions. Do you see any interesting patterns? Feel free to share them with myself and others in the class.